



Automatic handgun detection alarm in videos using deep learning



Roberto Olmos^a, Siham Tabik^{a,*}, Francisco Herrera^{a,b}

^a Soft Computing and Intelligent Information Systems research group, University of Granada, Granada 18071, Spain

^b Faculty of Computing and Information Technology, King Abdulaziz University (KAU) Jeddah, Saudi Arabia

ARTICLE INFO

Article history:

Received 27 February 2017

Revised 2 May 2017

Accepted 10 May 2017

Available online 15 May 2017

Communicated By Dr Hu Jun

Keywords:

Classification

Detection

Deep learning

Convolutional Neural Networks

Faster R-CNN

VGG-16

Alarm Activation Time per Interval

ABSTRACT

Current surveillance and control systems still require human supervision and intervention. This work presents a novel automatic handgun detection system in videos appropriate for both, surveillance and control purposes. We reformulate this detection problem into the problem of minimizing false positives and solve it by i) building the key training data-set guided by the results of a deep Convolutional Neural Networks (CNN) classifier and ii) assessing the best classification model under two approaches, the sliding window approach and region proposal approach. The most promising results are obtained by Faster R-CNN based model trained on our new database. The best detector shows a high potential even in low quality youtube videos and provides satisfactory results as automatic alarm system. Among 30 scenes, it successfully activates the alarm after five successive true positives in a time interval smaller than 0.2 s, in 27 scenes. We also define a new metric, Alarm Activation Time per Interval (AATpl), to assess the performance of a detection model as an automatic detection system in videos.

© 2017 Elsevier B.V. All rights reserved.

1. Introduction

The crime rates caused by guns are very concerning in many places in the world, especially in countries where the possession of guns is legal or was legal for a period of time. The last statistics reported by the United Nations Office on Drugs and Crime (UNODC) reveals that the number of crimes involving guns per 100,000 habitants is very high in many countries, e.g., 21.5 in Mexico, 4.7 in United States and 1.6 in Belgium [19]. In addition, several psychological studies demonstrated that the simple fact of having access to a gun increases drastically the probability of committing a violent behavior [25].

One way to reducing this kind of violence is prevention via early detection so that the security agents or policemen can act. In particular, one innovative solution to this problem is to equip surveillance or control cameras with an accurate automatic handgun detection alert system. Related studies address the detection of guns but only on X-ray or millimetric wave images and only using traditional machine learning methods [6,7,26,27,29].

In the last five years, deep learning in general and Convolutional Neural Networks (CNNs) in particular have achieved superior results to all the classical machine learning methods in image classification, detection and segmentation in several applica-

tions [8,13,18,22,23,30]. Classical methods require manual intervention, whereas deep CNNs models automatically discover increasingly higher level features from data [11,17]. We aim at developing a good gun detector in videos using CNNs.

A proper training of deep CNNs, which contain millions of parameters, requires very large datasets, in the order of millions of samples, as well as High Performance Computing (HPC) resources, e.g., multi-processor systems accelerated with GPUs. Transfer learning through fine-tuning is becoming a widely accepted alternative to overcome these constraints. It consists of reutilizing the knowledge learnt from one problem to another related one [20]. Applying transfer learning to deep CNNs depends on the similarities between the original and new problem and also on the size of the new training set.

In general, fine-tuning the entire network, i.e., updating all the weights, is only used when the new dataset is large enough, else the model could suffer overfitting especially among the first layers of the network. Since these layers extract low-level features, e.g., edges and color, they do not change significantly and can be reutilized for several visual recognition tasks. The last layers of the CNN are gradually adjusted to the particularities of the problem to extract high level features. In this work we used a VGG-16 based classification model pre-trained on the ImageNet dataset (around 1.28 million images over 1000 generic object classes) [24] and fine-tuned on our own dataset of 3000 images of guns taken in a variety of contexts.

* Corresponding author.

E-mail addresses: siham@ugr.es, siham.tabik@gmail.com (S. Tabik), herrera@decsai.ugr.es (F. Herrera).

Using CNNs to automatically detect pistols in videos faces several challenges:

- Pistols can be handled with one or two hands in different ways and thus a large part of the pistol can be occluded.
- The process of designing a new dataset is manual and time consuming.
- The labeled dataset cannot be re-utilized by different detection approaches since they require different preprocessing, different labeling operations and cannot learn from the same labeled databases.
- Automatic pistol detection alarm requires the activation of the alarm in real time and only when the system is confident about the existence of a pistol in the scene.
- Automatic detection alarm systems require an accurate location of the pistol in the monitored scene.

As far as we know, this work presents the first automatic gun detection alarm system that uses deep CNNs based detection models. We focus on the most used type of handguns in crimes [31], pistol, which includes, revolver, automatic and semi-automatic pistols, six-gun shooters, horse pistol and derringers. To guide the design of the new dataset and to find the best detector we consider the following steps:

- we reformulate the problem of automatic pistol detection alarm in videos into the problem of minimizing the number of false positives where *pistol* represents the true class and
- we evaluate and compare the VGG-16 based classifier using two different detection approaches, the sliding window and region proposals approaches.

Due to the particularities of each approach, we applied different optimizations in each case. We evaluated increasing the number of classes in the sliding window approach and designing a richer training dataset for the region proposals approach.

As this work focuses on near real time solutions, we selected the most accurate and fastest detector and assess its performance on seven videos of different characteristics. Then, we evaluated its suitability as automatic pistol detection alarm system using a new metric, the Alarm Activation Time per Interval (AATpl), that measures the activation time for each scene with guns.

The main contributions of this work are:

- Designing a new labeled database that makes the learning model achieve high detection qualities. Our experience in building the new dataset and detector can be useful to guide developing the solution of other different problems.
- Finding the most appropriate CNN-based detector that achieves real-time pistol detection in videos.
- Introducing a new metric, AATpl, to assess the suitability of the proposed detector as automatic detection alarm system.

From the experiments we found that the most promising results are obtained by Faster R-CNN based model trained on our new database. The best performing model shows a high potential even in low quality youtube videos and provides satisfactory results as automatic alarm system. Among 30 scenes, it successfully activates the alarm, after five successive true positives, within an interval of time smaller than 0.2 s, in 27 scenes.

This paper is organized as follows. Section 2 gives a brief analysis of the most related papers. Section 3 provides an overview of the CNN model used in this work. Section 4 describes the procedure we have used to find the best detector that reaches good precisions and low false positives rate. Section 5 analyzes the performance of the built detector using seven videos and introduces a new metric to assess the performance of the detector as automatic detection system. Finally the conclusions are summarized in Section 6.

2. Related works

The problem of handgun detection in videos using deep learning is related in part to two broad research areas. The first addresses gun detection using classical methods and the second focuses on improving the performance of object detection using deep CNNs.

2.1. Gun detection

The first and traditional sub-area in gun detection focuses on detecting concealed handguns in X-ray or millimetric wave images. The most representative application in this context is luggage control in airports. The existent methods achieve high accuracies by using different combinations of feature extractors and detectors, either using simple density descriptors [6], border detection and pattern matching [7] or using more complex methods such as cascade classifiers with boosting [29]. The effectiveness of these methods made them essential in some specific places. However, they have several limitations. As these systems are based on metal detection, they cannot detect non metallic guns. They are expensive to be used in many places as they require to be combined with X-ray scanners and Conveyor belts. They are not precise because they react to all metallic objects.

The second sub-area addresses gun detection in RGB images using classical methods. The few existent papers essentially apply methods such as SIFT (Scale-Invariant Feature Transform) y RIFT (Rotation-Invariant Feature Transform), combined with Harris interest point detector or FREAK (Fast Retina Keypoint) [12,26,27]. For example, the authors in [26,27] developed an accurate software for pistol detection in RGB images. However, their method is unable to detect multiple pistols in the same scene. The used technique consists of first, eliminating non related objects to a pistol from the segmented image using K-mean clustering algorithm then, applying SURF (Speeded Up Robust Features) method for detecting points of interest. Similarly, the authors in [12] demonstrated that BoWSS (Bag of Words Surveillance System) algorithm has a high potential to detect guns. They first extract features using SIFT, cluster the obtained functions using K-Means clustering and use SVM (Support Vector Machine) for the training. The authors in [14] addresses rifle detection in RGB images using SVM (Support Vector Machine).

All the above cited systems are slow, cannot be used for constant monitoring, require the supervision of an operator and cannot be used in open areas.

2.2. Detection models

Object detection consists of recognizing the object and finding its location in the input image. The existing methods address the detection problem by reformulating it into a classification problem, they first train the classifier then during the detection process they run it on a number of areas of the input image using either the sliding window approach or region proposals approach.

- **Sliding window** approach: It is an exhaustive method that considers a large number of candidate windows, in the order of 10^4 , from the input image. It scans the input image, at all locations and multiple scales, with a window and runs the classifier at each one of the windows. The most relevant works in this context improve the performance of the detection by building more sophisticated classifiers. The Histogram of Oriented Gradients (HOG) based model [3] uses HOG descriptor for feature extraction to predict the object class in each window. The Deformable Parts Models (DPM) [4], which is an extension of HOG based model, uses (1) HOG descriptor to calculate low-level

Table 4

The total number of True Positives #TP, total number of Ground Truth true Positives #GT_P, total number of False Positives #FP in the considered seven videos, labeled as video 1 to 7.

| video# | #frames | #TP | #GT_P | #FP | Precision (%) | Recall (%) | F1 measure (%) |
|--------|---------|-----|-------|-----|---------------|------------|----------------|
| 1 | 393 | 60 | 162 | 8 | 88.24 | 37.04 | 52.17 |
| 2 | 627 | 467 | 778 | 11 | 98.70 | 60.03 | 74.36 |
| 3 | 441 | 25 | 58 | 15 | 62.50 | 43.10 | 51.02 |
| 4 | 591 | 6 | 54 | 0 | 100.00 | 11.11 | 20.22 |
| 5 | 627 | 24 | 105 | 21 | 53.33 | 22.86 | 32.00 |
| 6 | 212 | 141 | 290 | 30 | 82.46 | 48.62 | 61.17 |
| 7 | 501 | 166 | 476 | 6 | 96.51 | 34.87 | 51.23 |

databases built in this work are summarized in Table 1. In the region proposals approach, the detector is directly trained on region proposals of a new database, Database-5, with richer contexts.

To evaluate and compare the performance of the classification model trained on the five proposed databases, we have built a test-set of 608 images distributed into 304 images that contain pistols and 304 images that do not contain pistols. We have used three metrics, *precision*, *recall* and *F1 measure*, which evaluates the balance between the *precision* and *recall*. Where

$$\text{precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

$$\text{recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

and

$$\text{F1 measure} = 2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$$

For implementing the proposed models we used theano with Keras as front-end [1,2] in the sliding window based approach shown in Section 4.1 and Caffe [16] in the region proposal approach shown in Section 4.2.1.

4.1. Sliding window approach

This section aims at finding the best combination of classifier and training dataset within the sliding window detection approach. We address the minimization of the number of false positives by building different training datasets and increasing the number of classes.

4.1.1. Two classes

To guide our design, and as a first approximation, we have built a preliminary database, Database-1, and consider a broader class of type gun. Database-1 contains 9100 images extracted from online gun catalogues, gun use tutorials and gun advertisement (see details in Table 1). We distributed it into 3990 images of class gun, which includes different types of guns, pistols, machine guns, rifle, machine guns, grenade, multiple rocket launcher, tank and 5110 images of not-gun class, which contains different kind of objects.

We have evaluated the VGG-16 based classification model considering two classes to indicate either the presence or absence of a pistol in the scene. Table 2 provides the number of true positives, #TP, the number of false positives, #FN, the number of true negatives, #TN, the number of false positives, #FP, precision, recall and F1 measure obtained by the classification model when trained on Database-1, 2, 3 and 4 respectively.

As shown in this Table, the classification model trained on Database-1 obtains a high number of false positives, 191, high number of false negatives, 272, and consequently low precision and recall. This can be explained by the fact that the large variety of guns makes the learning task very difficult. To make the problem affordable, we instead focus our task to pistol detection.



Fig. 2. Examples from Database-2, the top three images represent the pistol class and the down three images represent the background class.

In a second step, we have built Database-2 using 1857 images, 751 images of class pistol and 1056 of background. We included to the background class different images of hands holding different objects other than pistols, e.g., cell phone and pen as illustrated in Fig. 2. On the test set, the binomial classification model trained on Database-2 obtained 11 false positives, a high number of false negatives 206, a precision of 89.91%, recall 32.24% and F1 measure 47.46%, which are still below our expectations. By analyzing the false positives we found that most of them consider the white background as part of the pistol which is due to the presence of the white background in most training examples.

4.1.2. Multiple classes

To further decrease the number of false positives we designed two new databases, Database-3 and Database-4, and considered a higher number of classes.

Database-3 contains 5470 images set distributed into 28 classes. 751 of training examples contain pistols. As shown in Table 2, the 28-classes based classification model overcomes the binomial model by reducing the number of false positives to 5 and improving the precision to 94.44%, recall to 27.96% and F1 measure to 43.15%. However, the number of false negatives is still very high.

We have also explored the performance of the classification model using a higher number of classes on Database-4. Database-4 is built using 9261 images distributed into 102 classes, 200 training examples contain pistols and 9061 contain diverse objects, e.g., airplanes, ant, beaver, brain, chair, cell phone, animals and so on. This classifier produces the best results over the previous ones, a lower number of false negatives 207, slightly lower number of true negatives 298 and consequently better recall 31.91% and F1 measure 47.67%. Next we will evaluate the detection process using the 102-classes based classifier.

4.1.3. Detection process and final analysis

We selected the best classification model to evaluate it under the sliding window approach. The classifier is applied automatically, in steps of 60 × 60-pixels, to windows of 160 × 120-pixels of each input image to determine whether it contains a pistol. The

whole detection process takes 1.5 seconds in a 640×360 -pixels input image. Which is not acceptable for an automatic system such as the one considered in this work.

The detection model that makes use of the sliding window approach achieves few false positives and high precision, however, it obtains a low recall 35% and its execution time is not appropriate for online detection. In next section we will explore another alternative to further improve performance and speed of the detection process.

4.2. Region proposals approach

This section analyzes the performance of the detection under the region proposals approach and suitability for real time analysis.

4.2.1. Learning model based on Regions and labeled pistol bounding box

In this approach, we have used Faster Region based-CNN (Faster R-CNN) [21], which combines the selective search method with VGG-16 based classifier.

The design of a new training dataset for this approach is also manual and cannot re-use the databases from the previous approach. We have built Database-5 using 3000 images that contains pistols in different contexts and scenarios, downloaded from diverse web-sites. Fig. 2, as complementary material¹, provides three examples of Database-5. We considered a two class model and labeled the pistols by providing its localization, i.e., bounding box, in each individual training image. The rest of objects in the image are considered background. It is worth noting that the training step takes around 52 minutes on Nvidia GTX 980.

In general, as it can be seen from Table 3, Faster R-CNN trained on Database-5 obtains the highest performance over all the previously analyzed models. It provides the highest true positives number and the highest true negative numbers and consequently the highest recall 100% and F1 score 91.43%. However it produces more false positives, 57, and consequently lower precision, 84.21%. In Section 5, we address this issue in the context of automatic alarm systems by activating the alarm only when at least five successive false positives happen in five successive frames. Next we analyze the speed of this detection model.

4.2.2. Detection process and final analysis

We have evaluated the detection process using Faster-RCNN, which is based on the region proposals approach. We considered two classes and trained the classification model on Database 5. The whole detection process in a 1000×1000 -pixels image takes 0.19 s approximately producing a rate of 5.3 frames/s. Which allows the pistol detection to be performed in videos in near real time.

The detection model based on region proposals approach achieves the maximum recall 100% over the pistol class, zero false negatives, good precision 85.21%, a reasonable false positives number and can be used for real time detection. This makes it a good candidate for detecting pistols in a sequence of frames as shown in Section 5.

5. Analysis of the detection in videos

In this section we explore the strengths and weaknesses of our model on seven low quality youtube videos. In particular, we first assess the quality of the detection and localization (Section 5.1) then analyze the suitability of our model as pistol detection alarm system using a new metric (Section 5.2).

5.1. Analysis of detection and localization

This section analyzes the performance of our best detection model on seven public videos, with low qualities, downloaded from youtube. Six of them are pieces of well known films from the 90s, James Bon: The World is Not Enough (video 1), 3 pieces from Pulp Fiction (video 2, 3 and 4), Impossible Mission: Rogue Nation (video 5) and Mister Bean (video 6). The seventh video is a long pistol threatening video (video 7). The videos with the detections can be found in a public repository in github².

For the experiments, we analyzed the results of the detection in the videos frame by frame and consider a detection as true positive if the overlapping between the handled pistol and the predicted bounding box is more than 50%. Recall that the way pistols are handled is also a key to the detection. We consider a pistol as ground truth when it is recognizable by the human eye. Table 4 provides the total number of True Positives #TP, the total number of False Positives #FP and total number of Ground Truth Positives #GT_P, in each one of the seven videos. We consider a threshold $\in [0.70.9]$.

In general, although the scenes are dynamic in most videos, the detector achieves good balance between precision and recall, especially in videos 2 and 6. Fig. 3, as complementary material³, shows one example of an accurate detection. In particular, the detector provides very high precisions in videos 1, 2, 3, 4, 6 and 7, and the obtained number of false positives is very low in all the videos. Which is essential to avoid activating negative alarms. The obtained false positive detections can be addressed in a realistic system by activating the alarm only when false positives are detected in a number of consecutive frames.

The low recall can be explained by the false negatives detected in the frames with very low contrast and luminosity. Fig. 4, as complementary material⁴, shows two frames, of a gadget, with low quality. The false negatives depends on the quality of the frame and whether the pistol is clearly visible. In particular, they occur when the pistol is moved very fast or when it is placed in the background. Fig. 5, as complementary material, shows an example of two false negatives in the background.

In conclusion, the obtained results can be considered as acceptable to detect sequences of clearly visible pistols as it will be analyzed in next section.

5.2. Analyzing the model as alarm detection system

In an automatic pistol detection system the alarm must be activated when the system is completely confident about the presence of pistols in the scene. To assess the performance of our detection model as an alarm system, we define a new metric, *Alarm Activation Time per Interval* (AATpI). AATpI quantifies how fast is the system in detecting pistols in a given scene.

Definition 1. AATpI is the time the automatic detection alarm system takes to detect at least k successive frames of true positives.

In the next analysis we used $k = 5$. For the experiments, we selected 30 scenes from the previously used videos with the next requirements. Each scene is made up of at least 5 frames, filmed in a fixed scenario, i.e., in the same place, and the pistols are clearly

² <https://github.com/SihamTabik/Pistol-Detection-in-Videos.git>.

³ Figure 3 can be visualized through this public link: <https://github.com/SihamTabik/Pistol-Detection-in-Videos/blob/master/ComplementaryMaterial.pdf>.

⁴ Figs. 4 and 5 can be visualized them through this public link: <https://github.com/SihamTabik/Pistol-Detection-in-Videos/blob/master/ComplementaryMaterial.pdf>.

¹ Fig. 2 is not explicitly shown in this paper because it may contain violent connotations. The reader can visualize it through this public link: <https://github.com/SihamTabik/Pistol-Detection-in-Videos/blob/master/ComplementaryMaterial.pdf>.

visible to a human viewer. These scenes can be found in a public repository in github⁵.

The model successfully detects the pistol in 27 scenes with an average time interval AATpI=0.2 s, which is good enough for an alarm system. The detector fails to detect pistols only in three scenes. This is due to the same reasons highlighted previously, which are the low contrast and luminosity of the frames, the pistol is moved very fast or when the pistol is not in the foreground.

In summary, although we have used low quality videos for the evaluation, the proposed model has shown good performance and demonstrated to be appropriate for automatic pistol detection alarm systems.

6. Conclusions and future work

This work presented a novel automatic pistol detection system in videos appropriate for both, surveillance and control purposes. We reformulate this detection problem into the problem of minimizing false positives and solve it by building the key training data-set guided by the results of a VGG-16 based classifier, then assessing the best classification model under two approaches, the sliding window approach and region proposal approach. The most promising results have been obtained with Faster R-CNN based model, trained on our new database, providing zero false positives, 100% recall, a high number of true negatives and good precision 84,21%. The best detector has shown a high potential even in low quality youtube videos and provides very satisfactory results as automatic alarm system. Among 30 scenes, it successfully activates the alarm after five successive true positives within an interval of time smaller than 0.2 s, in 27 scenes.

The proposed detector can be used in several applications, e.g., i) real time detection of guns in places monitored by cameras and ii) control whether the videos uploaded to social media contain scenes with guns.

As present and future work, we are evaluating reducing the number of false positives, of Faster R-CNN based detector, by pre-processing the videos, i.e., increasing their contrast and luminosity, and also by enriching the training set with pistols in motion. We will also evaluate different CNNs-based classifier such as, GoogLeNet and consider a higher number of classes.

Acknowledgments

This work was partially supported by the Spanish Ministry of Science and Technology under the project TIN2014-57251-P. Siham Tabik was supported by the Ramón y Cajal Programme (RYC-2015-18136).

References

- [1] R. Al-Rfou, G. Alain, A. Almahairi, C. Angermueller, D. Bahdanau, N. Ballas, F. Bastien, J. Bayer, A. Belikov, A. Belopolsky, et al., Theano: A python framework for fast computation of mathematical expressions, arXiv preprint arXiv:1605.02688 (2016).
- [2] F. Chollet, Keras: Theano-based deep learning library, Code: <https://github.com/fchollet>. Documentation: <http://keras.io> (2015).
- [3] N. Dalal, B. Triggs, Histograms of oriented gradients for human detection, in: 2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'05), 1, IEEE, 2005, pp. 886–893.
- [4] P.F. Felzenszwalb, R.B. Girshick, D. McAllester, D. Ramanan, Object detection with discriminatively trained part-based models, IEEE Trans. Pattern Anal. Mach. Intell. 32 (9) (2010) 1627–1645.
- [5] P.F. Felzenszwalb, D.P. Huttenlocher, Pictorial structures for object recognition, Int. J. Comput. Vis. 61 (1) (2005) 55–79.
- [6] G. Flitton, T.P. Breckon, N. Megherbi, A comparison of 3d interest point descriptors with application to airport baggage object detection in complex ct imagery, Pattern Recogn. 46 (9) (2013) 2420–2436.

- [7] R. Gesick, C. Saritac, C.-C. Hung, Automatic image analysis process for the detection of concealed weapons, in: Proceedings of the 5th Annual Workshop on Cyber Security and Information Intelligence Research: Cyber Security and Information Intelligence Challenges and Strategies, ACM, 2009, p. 20.
- [8] M.M. Ghazi, B. Yanikoglu, E. Aptoula, Plant identification using deep neural networks via optimization of transfer learning parameters, Neurocomputing 235 (2017) 228–235.
- [9] R. Girshick, Fast r-cnn, in: Proceedings of the IEEE International Conference on Computer Vision, 2015, pp. 1440–1448.
- [10] R. Girshick, J. Donahue, T. Darrell, J. Malik, Rich feature hierarchies for accurate object detection and semantic segmentation, in: Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2014, pp. 580–587.
- [11] Y. Guo, Y. Liu, A. Oerlemans, S. Lao, S. Wu, M.S. Lew, Deep learning for visual understanding: a review, Neurocomputing 187 (2016) 27–48.
- [12] N.B. Halima, O. Hosam, Bag of words based surveillance system using support vector machines, Int. J. Secur. Appl. 10 (4) (2016) 331–346.
- [13] G. Hinton, L. Deng, D. Yu, G. Dahl, A.-r. Mohamed, N. Jaitly, A. Senior, V. Vanhoucke, P. Nguyen, T. Sainath, B. Kingsbury, Deep neural networks for acoustic modeling in speech recognition: the shared views of four research groups, IEEE Signal Process. Mag. 29 (6) (2012) 82–97.
- [14] O. Hosam, A. Alraddadi, K-means clustering and support vector machines approach for detecting fire weapons in cluttered scenes, Life Sci. J. 11 (9) (2014).
- [15] J. Hosang, R. Benenson, P. Dollár, B. Schiele, What makes for effective detection proposals? IEEE Trans. Pattern Anal. Mach. Intell. 38 (4) (2016) 814–830.
- [16] Y. Jia, E. Shelhamer, J. Donahue, S. Karayev, J. Long, R. Girshick, S. Guadarrama, T. Darrell, Caffe: Convolutional architecture for fast feature embedding, arXiv preprint arXiv:1408.5093 (2014).
- [17] A. Krizhevsky, I. Sutskever, G.E. Hinton, Imagenet classification with deep convolutional neural networks, in: Advances in Neural Information Processing Systems, 2012, pp. 1097–1105.
- [18] Q.V. Le, Building high-level features using large scale unsupervised learning, in: 2013 IEEE International Conference Acoustics Speech Signal Processing, 2013, pp. 8595–8598.
- [19] United Nations Office on Drugs and Crime (UNODC), Global study on homicide 2013, Data: UNODC Homicide Statistics 2013(2013).
- [20] S.J. Pan, Q. Yang, A survey on transfer learning, IEEE Trans. Knowl. Data Eng. 22 (10) (2010) 1345–1359.
- [21] S. Ren, K. He, R. Girshick, J. Sun, Faster r-cnn: towards real-time object detection with region proposal networks, in: Advances in Neural Information Processing Systems, 2015, pp. 91–99.
- [22] T.N. Sainath, A.-r. Mohamed, B. Kingsbury, B. Ramabhadran, Deep convolutional neural networks for LVCSR, in: 2013 IEEE International Conference Acoustics Speech Signal Processing, 2013, pp. 8614–8618.
- [23] X. Shu, Y. Cai, L. Yang, L. Zhang, J. Tang, Computational face reader based on facial attribute estimation, Neurocomputing 236 (2017) 153–163.
- [24] K. Simonyan, A. Zisserman, Very deep convolutional networks for large-scale image recognition, arXiv preprint arXiv:1409.1556(2014).
- [25] J.W. Swanson, N.A. Sampson, M.V. Petukhova, A.M. Zaslavsky, P.S. Appelbaum, M.S. Swartz, R.C. Kessler, Guns, impulsive angry behavior, and mental disorders: results from the national comorbidity survey replication (ncs-r), Behav. Sci. Law 33 (2–3) (2015) 199–212.
- [26] R.K. Tiwari, G.K. Verma, A computer vision based framework for visual gun detection using harris interest point detector, Proc. Comput. Sci. 54 (2015a) 703–712.
- [27] R.K. Tiwari, G.K. Verma, A computer vision based framework for visual gun detection using surf, in: Electrical, Electronics, Signals, Communication and Optimization (EESCO), 2015 International Conference on, IEEE, 2015b, pp. 1–5.
- [28] J.R. Uijlings, K.E. Van De Sande, T. Gevers, A.W. Smeulders, Selective search for object recognition, Int. J. Comput. Vis. 104 (2) (2013) 154–171.
- [29] Z. Xiao, X. Lu, J. Yan, L. Wu, L. Ren, Automatic detection of concealed pistols using passive millimeter wave imaging, in: 2015 IEEE International Conference on Imaging Systems and Techniques (IST), IEEE, 2015, pp. 1–4.
- [30] W. Yu, K. Yang, H. Yao, X. Sun, P. Xu, Exploiting the complementary strengths of multi-layer cnn features for image retrieval, Neurocomputing 237 (2017) 235–241.
- [31] M.W. Zawitz, in: Guns used in crime, Washington, DC: US Department of Justice: Bureau of Justice Statistics Selected Findings, publication NCJ-148201, 1995.



Roberto Olmos received the M.Sc. degree in Computer Science in 2015 from the Meritorious Autonomous University of Puebla, Puebla, Mexico. He is currently a Ph.D. student in the Department of computer Science and Artificial Intelligence, University of Granada. His research interests include object detection and supervised deep learning.

⁵ <https://github.com/SihamTabik/Pistol-Detection-in-Videos.git>.



Siham Tabik received the B.Sc. degree in physics from University Mohammed V, Rabat, Morocco, in 1998 and the Ph.D. degree in Computer Science from the University of Almería, Almería, Spain. She is currently Ramón y Cajal researcher at the University of Granada. Her research interests include the design of scalable algorithms, deep learning CNNs models and object detection.



Francisco Herrera received his M.Sc. in Mathematics in 1988 and Ph.D. in Mathematics in 1991, both from the University of Granada, Spain. He is currently a Professor in the Department of Computer Science and Artificial Intelligence at the University of Granada.

He has been the supervisor of 40 Ph.D. students. He has published more than 300 journal papers that have received more than 51000 citations (Scholar Google, H-index 114). He is coauthor of the books "Genetic Fuzzy Systems" (World Scientific, 2001) and "Data Preprocessing in Data Mining" (Springer, 2015), "The 2-tuple Linguistic Model. Computing with Words in Decision Making" (Springer, 2015), "Multilabel Classification. Problem analysis, metrics and techniques" (Springer, 2016), "Multiple Instance Learning. Foundations and Algorithms" (Springer, 2016).

He currently acts as Editor in Chief of the international journals "Information Fusion" (Elsevier) and "Progress in Artificial Intelligence" (Springer). He acts as editorial member of a dozen of journals.

He received the following honors and awards: ECCAI Fellow 2009, IFSA Fellow 2013, 2010 Spanish National Award on Computer Science ARITMEL to the "Spanish Engineer on Computer Science", International Cajastur "Mamdani" Prize for Soft Computing (Fourth Edition, 2010), IEEE Transactions on Fuzzy Systems Outstanding 2008 and 2012 Paper Award (bestowed in 2011 and 2015 respectively), 2011 Lotfi A. Zadeh Prize Best paper Award of the International Fuzzy Systems Association, 2013 AEPIA Award to a scientific career in Artificial Intelligence, and 2014 XV Andalucía Research Prize Maimónides and 2017 Andalucía Medal (by the regional government of Andalucía). He has been selected as a Thomson Reuters Highly Cited Researcher <http://highlycited.com/> (in the fields of Computer Science and Engineering, respectively, 2014 to present.)

His current research interests include among others, soft computing (including fuzzy modeling and evolutionary algorithms), information fusion, decision making, biometric, data preprocessing, data science and big data.